Hyperparameter tuning, Batch Normalization, Programming Frameworks

测验, 10 个问题

第 1 个问题

1

point

1。第 1 个问题

If searching among a large number of hyperparameters, you should try values in a grid rather than random values, so that you can carry out the search more systematically and not rely on chance. True or False?

True

False

第 2 个问题

1

point

2。第 2 个问题

Every hyperparameter, if set poorly, can have a huge negative impact on training, and so all hyperparameters are about equally important to tune well. True or False?

True

False

第 3 个问题

1

point

3。第 3 个问题

During hyperparameter search, whether you try to babysit one model ("Panda" strategy) or train a lot of models in parallel ("Caviar") is largely determined by:

Whether you use batch or mini-batch optimization

The presence of local minima (and saddle points) in your neural network

The amount of computational power you can access

The number of hyperparameters you have to tune

第 4 个问题

1

point

4。第 4 个问题

If you think β (hyperparameter for momentum) is between on 0.9 and 0.99, which of the following is the recommended way to sample a value for beta?

```
1  r = np.random.rand()
2  beta = r*0.09 + 0.9
```

```
1  r = np.random.rand()
2  beta = 1-10**(- r - 1)
```

```
1  r = np.random.rand()
2  beta = 1-10**(- r + 1)
```

```
1  r = np.random.rand()
2  beta = r*0.9 + 0.09
```

第 5 个问题
1
point

5。第 5 个问题

Finding good hyperparameter values is very time-consuming. So typically you should do it once at the start of the project, and try to find very good hyperparameters so that you don't ever have to revisit tuning them again. True or false?

True

False

第 6 个问题
1
point

6。第 6 个问题

In batch normalization as presented in the videos, if you apply it on the lth layer of your neural network, what are you normalizing?

b[l]

W[l]

a[l]

z[l]

第 7 个问题

7。第 7 个问题

In the normalization formula $z_{norm}^{(i)} = \frac{z^{(i)} - \mu}{\sqrt{\sigma^2 + \varepsilon}}$, why do we use epsilon?

To avoid division by zero

In case μ is too small

To have a more accurate normalization

To speed up convergence

8。第 8 个问题

Which of the following statements about γ and β in Batch Norm are true?

The optimal values are $\gamma = \sqrt{\sigma^2 + \varepsilon}$, and $\beta = \mu$.

There is one global value of $\gamma \in \Re$ and one global value of $\beta \in \Re$ for each layer, and applies to all the hidden units in that layer.

They can be learned using Adam, Gradient descent with momentum, or RMSprop, not just with gradient descent.

β and γ are hyperparameters of the algorithm, which we tune via random sampling.

They set the mean and variance of the linear variable $z^{[l]}$ of a given layer.

9。第 9 个问题

After training a neural network with Batch Norm, at test time, to evaluate the neural network on a new example you should:

Use the most recent mini-batch's value of μ and σ^2 to perform the needed normalizations.

Perform the needed normalizations, use μ and σ^2 estimated using an exponentially weighted average across mini-batches seen during training.

Skip the step where you normalize using μ and σ^2 since a single test example cannot be normalized.

If you implemented Batch Norm on mini-batches of (say) 256 examples, then to evaluate on one test example, duplicate that example 256 times so that you're working with a mini-batch the same size as during training.

第 10 个问题
1
point

10。第 10 个问题

Which of these statements about deep learning programming frameworks are true? (Check all that apply)

Even if a project is currently open source, good governance of the project helps ensure that the it remains open even in the long term, rather than become closed or modified to benefit only one company.

Deep learning programming frameworks require cloud-based machines to run.

A programming framework allows you to code up deep learning algorithms with typically fewer lines of code than a lower-level language such as Python.

我了解不是我自己完成的作业将永远不会通过该课程且我的 Coursera 帐号会被取消激活。